

P8106 HW1

Lin Yang

```
library(tidyverse)
library(corrplot)
library(leaps)
library(glmnet)
library(plotmo)
library(caret)
library(pls)
```

Import the training data and test data

```
train <- read.csv("data/housing_training.csv") %>%
  janitor::clean_names()
train <- na.omit(train)

test <- read.csv("data/housing_test.csv") %>%
  janitor::clean_names()
test <- na.omit(test)
```

Least squares

We first fit a linear model on the training data using least squares and cross-validation.

```
ctrl <- trainControl(method = "repeatedcv", number = 10, repeats = 5)
x <- model.matrix(sale_price ~ ., train)[, -1]
y <- train$sale_price

set.seed(1234)
fit_lm <- train(x, y,
               method = "lm",
               trControl = ctrl)
summary(fit_lm)
```

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -89864 -12424      416  12143 140205
```

```
##
## Coefficients: (1 not defined because of singularities)
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -4.985e+06  3.035e+06  -1.642  0.10076
## gr_liv_area      2.458e+01  1.393e+01   1.765  0.07778 .
## first_flr_sf     4.252e+01  1.409e+01   3.017  0.00260 **
## second_flr_sf    4.177e+01  1.379e+01   3.029  0.00250 **
## total_bsmt_sf    3.519e+01  2.744e+00  12.827 < 2e-16 ***
## low_qual_fin_sf      NA         NA      NA      NA
## wood_deck_sf      1.202e+01  4.861e+00   2.474  0.01350 *
## open_porch_sf     1.618e+01  1.004e+01   1.611  0.10736
## bsmt_unf_sf      -2.087e+01  1.723e+00 -12.116 < 2e-16 ***
## mas_vnr_area      1.046e+01  4.229e+00   2.473  0.01353 *
## garage_cars       4.229e+03  1.893e+03   2.234  0.02563 *
## garage_area       7.769e+00  6.497e+00   1.196  0.23195
## year_built        3.251e+02  3.130e+01  10.388 < 2e-16 ***
## tot_rms_abv_grd   -3.838e+03  6.922e+02  -5.545 3.51e-08 ***
## full_bath        -4.341e+03  1.655e+03  -2.622  0.00883 **
## overall_qualAverage -5.013e+03  1.735e+03  -2.890  0.00391 **
## overall_qualBelow_Average -1.280e+04  2.677e+03  -4.782 1.92e-06 ***
## overall_qualExcellent 7.261e+04  5.381e+03  13.494 < 2e-16 ***
## overall_qualFair    -1.115e+04  5.240e+03  -2.127  0.03356 *
## overall_qualGood     1.226e+04  1.950e+03   6.287 4.30e-10 ***
## overall_qualVery_Excellent 1.304e+05  8.803e+03  14.810 < 2e-16 ***
## overall_qualVery_Good 3.798e+04  2.741e+03  13.852 < 2e-16 ***
## kitchen_qualFair    -2.663e+04  6.325e+03  -4.210 2.71e-05 ***
## kitchen_qualGood    -1.879e+04  4.100e+03  -4.582 5.01e-06 ***
## kitchen_qualTypical -2.677e+04  4.281e+03  -6.252 5.37e-10 ***
## fireplaces         1.138e+04  2.257e+03   5.043 5.18e-07 ***
## fireplace_quFair    -7.207e+03  6.823e+03  -1.056  0.29106
## fireplace_quGood     6.070e+02  5.833e+03   0.104  0.91713
## fireplace_quNo_Fireplace 3.394e+03  6.298e+03   0.539  0.59002
## fireplace_quPoor    -5.185e+03  7.399e+03  -0.701  0.48362
## fireplace_quTypical -6.398e+03  5.897e+03  -1.085  0.27814
## exter_qualFair     -3.854e+04  8.383e+03  -4.598 4.66e-06 ***
## exter_qualGood     -1.994e+04  5.585e+03  -3.569  0.00037 ***
## exter_qualTypical   -2.436e+04  5.874e+03  -4.147 3.57e-05 ***
## lot_frontage       1.024e+02  1.905e+01   5.376 8.90e-08 ***
## lot_area           6.042e-01  7.864e-02   7.683 2.91e-14 ***
## longitude          -3.481e+04  2.537e+04  -1.372  0.17016
## latitude           5.874e+04  3.483e+04   1.686  0.09193 .
## misc_val           9.171e-01  1.003e+00   0.914  0.36071
## year_sold          -6.455e+02  4.606e+02  -1.401  0.16132
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 22190 on 1401 degrees of freedom
## Multiple R-squared:  0.9116, Adjusted R-squared:  0.9092
## F-statistic: 380.3 on 38 and 1401 DF, p-value: < 2.2e-16
```

```
#correlation plot
#corrplot(cor(x),
#          #method = "circle",
#          #type = "full",
```

```
#tl.cex = 0.5)
```

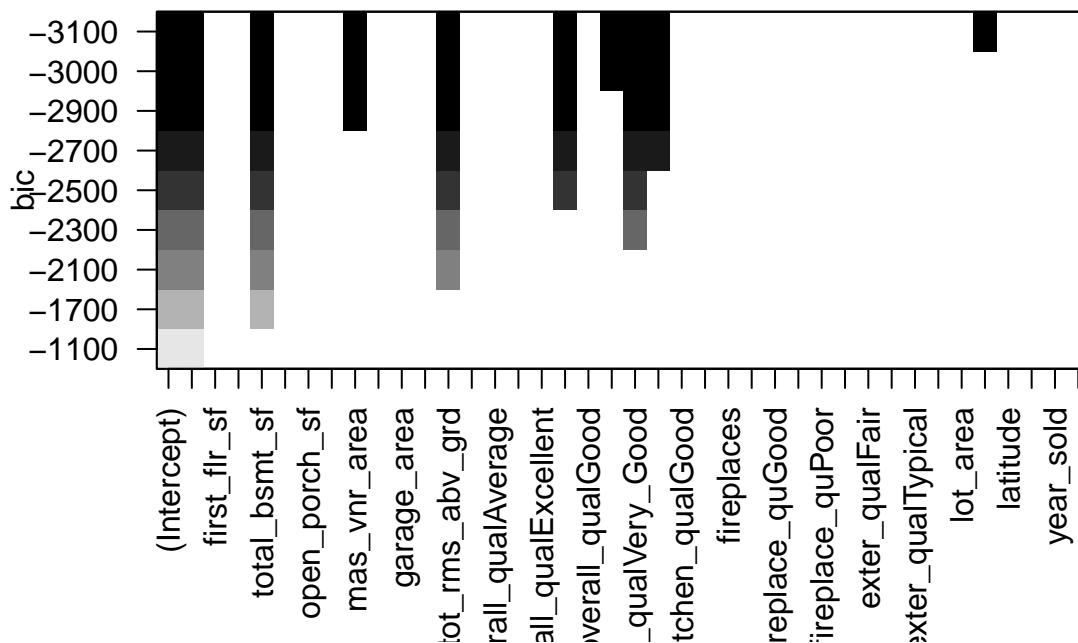
The least squares linear model is easy to fit, and the least squares estimates are BLUE. However, correlations amongst predictors can cause problems. From the above correlation plot, we can see that some predictors are highly correlated with each other, for example, `garage_area` and `garage_cars`. Due to multicollinearity, the variance of coefficients tends to increase and interpretations would be difficult.

We then did a best subset model selection, the predictors selected to give the smallest BIC are `gr_liv_area`, `total_bsmt_sf`, `mas_vnr_area`, `tot_rms_abv_grd`, `overall_qualFair`, `overall_qualVery_Excellent`, `overall_qualVery_Good`, `kitcehn_qualFair` and `longitude`.

```
regsubsetsObj <- regsubsets(sale_price ~ .,
                           data = train,
                           method = "exhaustive", nbest = 1)
```

```
## Reordering variables and trying again:
```

```
plot(regsubsetsObj, scale = "bic")
```

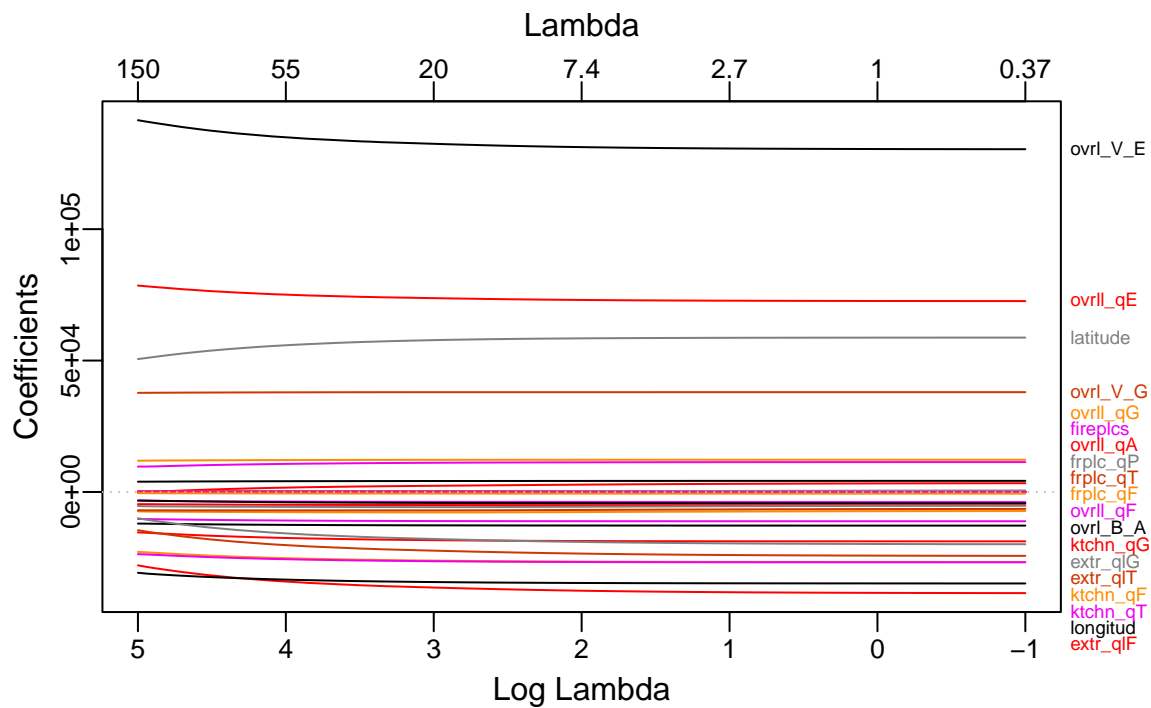


Lasso

Fit lasso model

```
fit_lasso <- glmnet(x = x,
                    y = y,
                    standardize = TRUE,
                    alpha = 1,
                    lambda = exp(seq(5, -1, length = 100)))

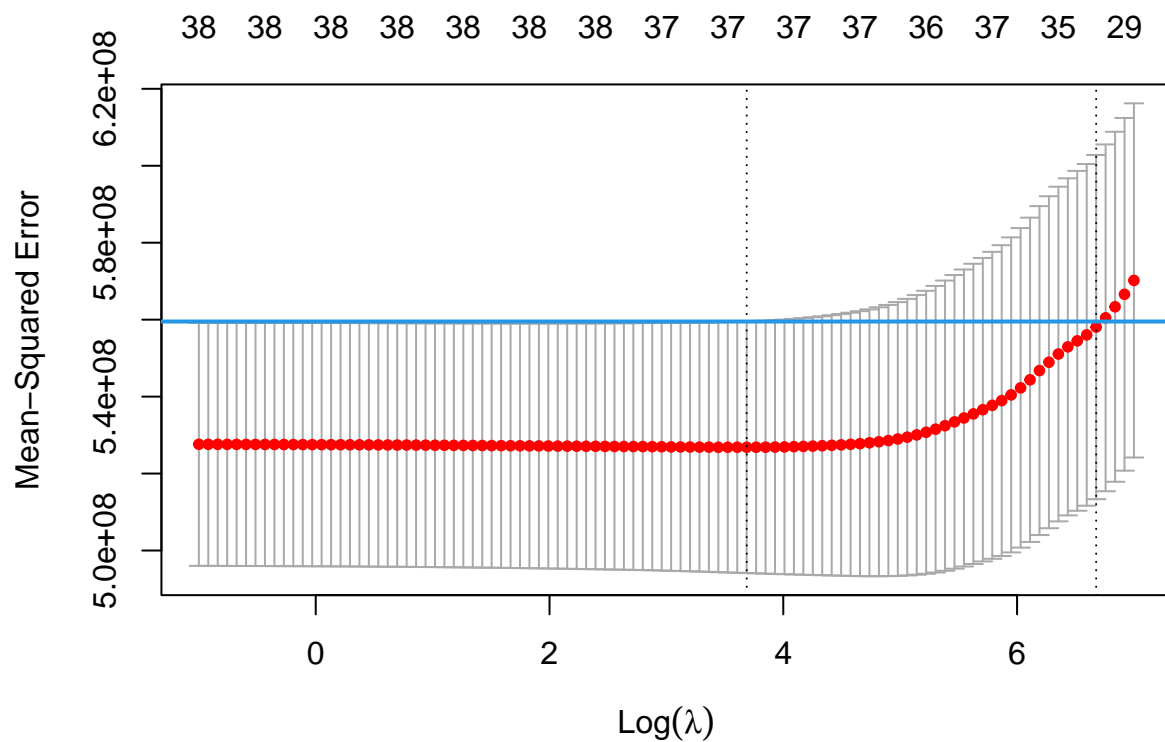
plot_glmnet(fit_lasso, xvar = "rlambda", label = 19)
```



Cross-validation for lasso

```
set.seed(1234)
cv.lasso <- cv.glmnet(x, y,
                      alpha = 1,
                      lambda = exp(seq(7, -1, length = 100)))

plot(cv.lasso) #cv curve
abline(h = (cv.lasso$cvm + cv.lasso$cvstd)[which.min(cv.lasso$cvm)], col = 4, lwd = 2)
```



```
#min CV MSE
cv.lasso$lambda.min
```

```
## [1] 39.91965
```

```
#1SE rule
cv.lasso$lambda.1se
```

```
## [1] 793.7493
```

```
#make prediction
x_test <- model.matrix(sale_price ~ ., test)[ , -1]
y_test <- test$sale_price
lasso_pred_min <- predict(cv.lasso, newx = x_test, s = "lambda.min", type = "response")
lasso_pred_1se <- predict(cv.lasso, newx = x_test, s = "lambda.1se", type = "response")
#test error
mean((lasso_pred_min - y_test)^2) #min MSE
```

```
## [1] 442495708
```

```
mean((lasso_pred_1se - y_test)^2) #1SE rule
```

```
## [1] 420534090
```

By performing cross-validation for the lasso model, the lambda with the minimal MSE is 39.9196501, and the lambda with 1SE rule is 793.7493067. The model with lambda.min gives a test error, 4.4249571×10^8 , and the model with lambda.1se gives a test error, 4.2053409×10^8 which is smaller, so 1 SE rule may be applied in this model.

Coefficients of the final lasso model

```
coef = predict(cv.lasso, s = cv.lasso$lambda.1se, type = "coefficients")
coef
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept)                  -2.387913e+06
## gr_liv_area                   5.693158e+01
## first_flr_sf                  1.089929e+00
## second_flr_sf                 .
## total_bsmt_sf                 3.679021e+01
## low_qual_fin_sf              -2.667815e+01
## wood_deck_sf                 8.375557e+00
## open_porch_sf                8.115431e+00
## bsmt_unf_sf                  -1.959769e+01
## mas_vnr_area                 1.418912e+01
## garage_cars                   3.077660e+03
## garage_area                  1.124366e+01
## year_built                   3.128514e+02
## tot_rms_abv_grd              -1.336851e+03
## full_bath                    .
## overall_qualAverage          -3.101488e+03
## overall_qualBelow_Average    -9.180078e+03
## overall_qualExcellent        9.044626e+04
## overall_qualFair             -6.412739e+03
## overall_qualGood             9.924709e+03
## overall_qualVery_Excellent   1.603936e+05
## overall_qualVery_Good        3.622543e+04
## kitchen_qualFair            -5.332043e+03
## kitchen_qualGood            .
## kitchen_qualTypical          -9.615242e+03
## fireplaces                   6.451761e+03
## fireplace_quFair            .
## fireplace_quGood            4.724947e+03
## fireplace_quNo_Fireplace     .
## fireplace_quPoor            .
## fireplace_quTypical          -1.043223e+02
## exter_qualFair              -1.465104e+04
## exter_qualGood              .
## exter_qualTypical           -5.086529e+03
## lot_frontage                 7.136431e+01
## lot_area                     5.618876e-01
## longitude                    -1.129777e+04
## latitude                     1.806072e+04
## misc_val                     .
## year_sold                    .
```

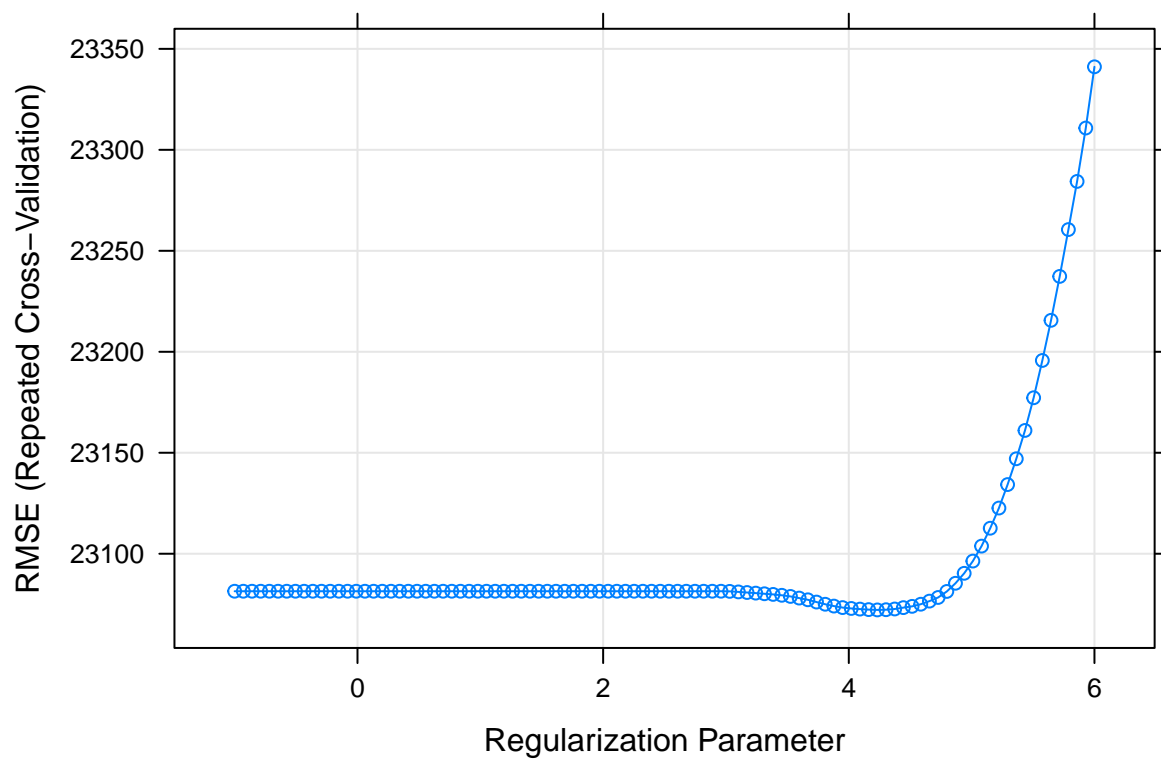
```
num_pred = length(which(coef != 0)) - 1
num_pred
```

```
## [1] 30
```

When 1SE rule is applied, 30 predictors are included in this model. The coefficients of predictors are shown above.

Lasso using caret.

```
set.seed(1234)
lasso.fit <- train(x, y,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = 1,
    lambda = exp(seq(6, -1, length = 100))),
  trControl = ctrl)
plot(lasso.fit, xTrans = log)
```



```
#optimal tuning parameters
lasso.fit$bestTune
```

```
##      alpha  lambda
## 75      1 68.87706
```

```
#coefficients
coef(lasso.fit$finalModel, lasso.fit$bestTune$lambda)
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)    -4.819134e+06
## gr_liv_area      6.533715e+01
## first_flr_sf      8.047852e-01
## second_flr_sf      .
## total_bsmt_sf     3.542815e+01
## low_qual_fin_sf   -4.088896e+01
## wood_deck_sf      1.161446e+01
## open_porch_sf     1.539102e+01
## bsmt_unf_sf       -2.088676e+01
## mas_vnr_area      1.092232e+01
## garage_cars       4.076890e+03
## garage_area       8.186225e+00
## year_built        3.232282e+02
## tot_rms_abv_grd   -3.605069e+03
## full_bath         -3.815672e+03
## overall_qualAverage -4.843987e+03
## overall_qualBelow_Average -1.243832e+04
## overall_qualExcellent 7.562416e+04
## overall_qualFair   -1.072979e+04
## overall_qualGood    1.211203e+04
## overall_qualVery_Excellent 1.359412e+05
## overall_qualVery_Good 3.788423e+04
## kitchen_qualFair   -2.474792e+04
## kitchen_qualGood   -1.711939e+04
## kitchen_qualTypical -2.523695e+04
## fireplaces         1.050261e+04
## fireplace_quFair   -7.656107e+03
## fireplace_quGood    .
## fireplace_quNo_Fireplace 1.371432e+03
## fireplace_quPoor   -5.631008e+03
## fireplace_quTypical -7.009951e+03
## exter_qualFair     -3.311445e+04
## exter_qualGood     -1.488441e+04
## exter_qualTypical  -1.932398e+04
## lot_frontage       9.950053e+01
## lot_area           6.042233e-01
## longitude          -3.283930e+04
## latitude            5.488323e+04
## misc_val            8.230881e-01
## year_sold          -5.559416e+02
```

Elastic net

Fit elastic net model


```

set.seed(1234)
fit_enet <- train(x, y,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
    lambda = exp(seq(8, -2, length = 50))),
  trControl = ctrl)
#best tuning parameters
fit_enet$bestTune

```

```

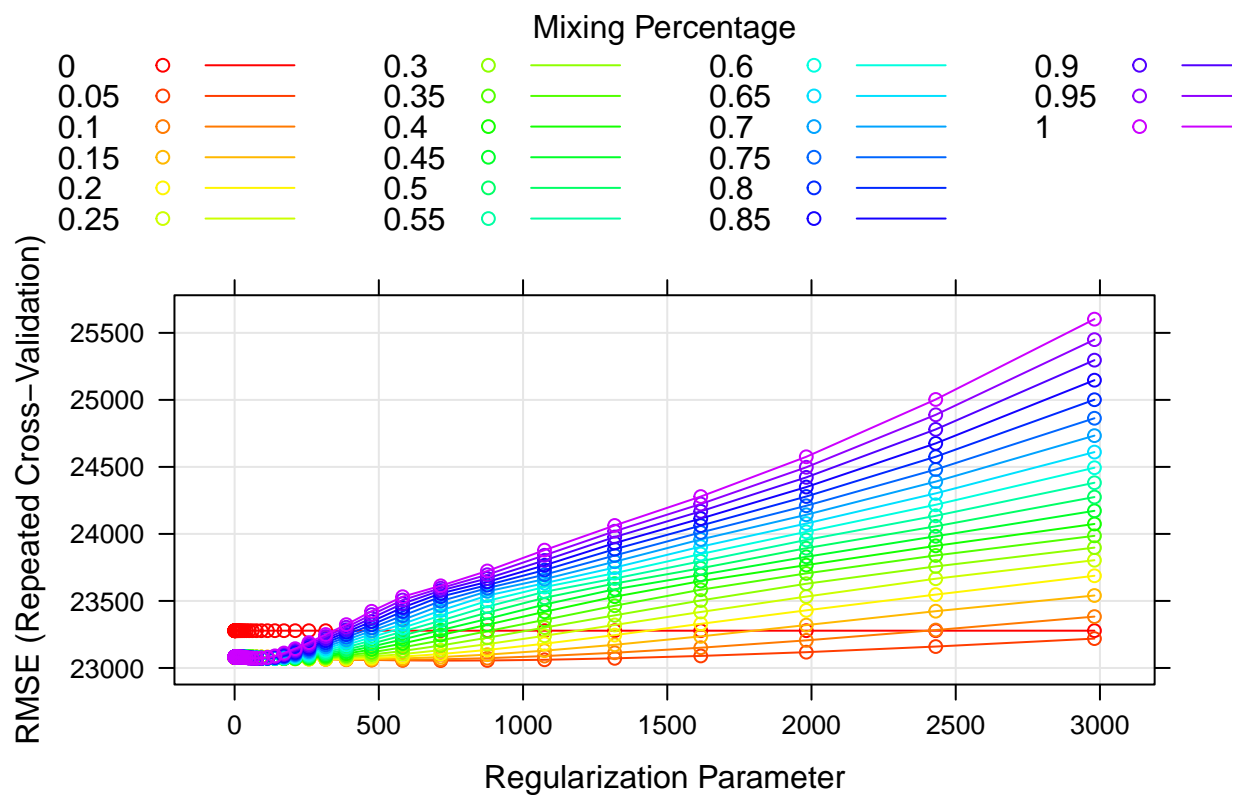
##      alpha      lambda
## 93  0.05 714.3897

```

```

myCol <- rainbow(25)
myPar <- list(superpose.symbol = list(col = myCol),
  superpose.line = list(col = myCol))
#plot of RMSE vs lambda
plot(fit_enet, par.settings = myPar)

```



The optimal tuning parameters are selected to be $\alpha = 0.05$ and $\lambda = 714.38967118394$.

Make predictions

```
enet_pred <- predict(fit_enet, newdata = x_test)
#test error
mean((enet_pred - y_test)^2)
```

```
## [1] 437161231
```

```
coef(fit_enet$finalModel, fit_enet$bestTune$lambda)
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept)                -5.143330e+06
## gr_liv_area                  3.857457e+01
## first_flr_sf                 2.660581e+01
## second_flr_sf                2.524578e+01
## total_bsmt_sf                3.488416e+01
## low_qual_fin_sf              -1.593502e+01
## wood_deck_sf                 1.239833e+01
## open_porch_sf                1.702473e+01
## bsmt_unf_sf                  -2.068247e+01
## mas_vnr_area                 1.189557e+01
## garage_cars                  4.019474e+03
## garage_area                  9.111479e+00
## year_built                   3.177757e+02
## tot_rms_abv_grd              -3.362468e+03
## full_bath                    -3.577411e+03
## overall_qualAverage          -5.132906e+03
## overall_qualBelow_Average    -1.268017e+04
## overall_qualExcellent         7.624393e+04
## overall_qualFair              -1.152272e+04
## overall_qualGood              1.190041e+04
## overall_qualVery_Excellent    1.372152e+05
## overall_qualVery_Good         3.753924e+04
## kitchen_qualFair              -2.314523e+04
## kitchen_qualGood              -1.561811e+04
## kitchen_qualTypical           -2.367917e+04
## fireplaces                    1.070460e+04
## fireplace_quFair              -7.908657e+03
## fireplace_quGood              1.177879e+02
## fireplace_quNo_Fireplace      1.556979e+03
## fireplace_quPoor              -5.851370e+03
## fireplace_quTypical           -6.995778e+03
## exter_qualFair                -3.215988e+04
## exter_qualGood                -1.376248e+04
## exter_qualTypical             -1.842191e+04
## lot_frontage                  9.967381e+01
## lot_area                      6.025529e-01
## longitude                     -3.529057e+04
## latitude                      5.758090e+04
## misc_val                      8.566244e-01
## year_sold                     -5.613626e+02
```

The test error of this elastic net model is 4.3716123×10^8 . The coefficients of this model are shown above.

Partial least squares

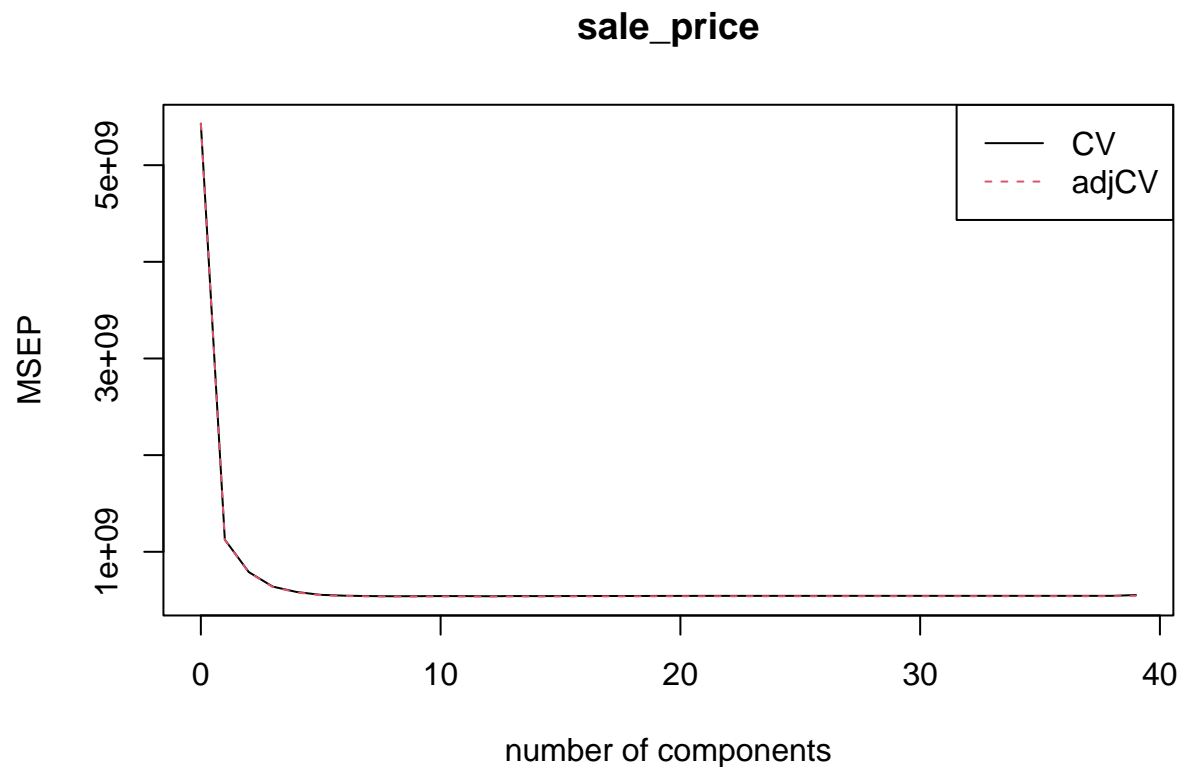
Fit partial least squares mode

```
set.seed(1234)
fit_pls <- plsr(sale_price ~ .,
               data = train,
               scale = TRUE,
               validation = "CV")
summary(fit_pls)
```

```
## Data:      X dimension: 1440 39
## Y dimension: 1440 1
## Fit method: kernelppls
## Number of components considered: 39
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              73685   33553   28106   25289   24162   23546   23362
## adjCV           73685   33537   28060   25207   24086   23471   23295
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV      23277   23238   23250   23272   23269   23240   23282
## adjCV    23210   23173   23182   23200   23196   23170   23207
##     14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## CV      23266   23279   23295   23290   23294   23312   23315
## adjCV    23193   23205   23219   23215   23219   23235   23238
##     21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
## CV      23323   23322   23322   23322   23323   23324   23326
## adjCV    23245   23245   23244   23244   23245   23246   23248
##     28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
## CV      23326   23326   23326   23327   23327   23327   23327
## adjCV    23248   23248   23248   23248   23248   23248   23248
##     35 comps 36 comps 37 comps 38 comps 39 comps
## CV      23327   23327   23327   23327   23506
## adjCV    23248   23248   23248   23248   23335
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X           20.02   25.93   29.67   33.59   37.01   40.03   42.49
## sale_price   79.73   86.35   89.36   90.37   90.87   90.99   91.06
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X           45.53   47.97   50.15   52.01   53.69   55.35   56.86
## sale_price   91.08   91.10   91.13   91.15   91.15   91.16   91.16
##     15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## X           58.64   60.01   62.18   63.87   65.26   67.10
## sale_price   91.16   91.16   91.16   91.16   91.16   91.16
##     21 comps 22 comps 23 comps 24 comps 25 comps 26 comps
## X           68.44   70.12   71.72   73.35   75.20   77.27
## sale_price   91.16   91.16   91.16   91.16   91.16   91.16
##     27 comps 28 comps 29 comps 30 comps 31 comps 32 comps
## X           78.97   80.10   81.83   83.55   84.39   86.34
## sale_price   91.16   91.16   91.16   91.16   91.16   91.16
```

```
##          33 comps  34 comps  35 comps  36 comps  37 comps  38 comps
## X          88.63    90.79    92.79    95.45    97.49    100.00
## sale_price    91.16    91.16    91.16    91.16    91.16    91.16
##          39 comps
## X          100.67
## sale_price    91.16
```

```
#plot of MSEP vs number of components
validationplot(fit_pls, val.type = "MSEP", legendpos = "topright")
```



```
#rmse of prediction
cv.mse <- RMSEP(fit_pls)
#number of components with the least rmsep
ncomp.cv <- which.min(cv.mse$val[1,,]) - 1
ncomp.cv
```

```
## 8 comps
##      8
```

```
#make predictions
pls_pred <- predict(fit_pls, newdata = x_test, ncomp = ncomp.cv)
#test error
mean((pls_pred - y_test)^2)
```

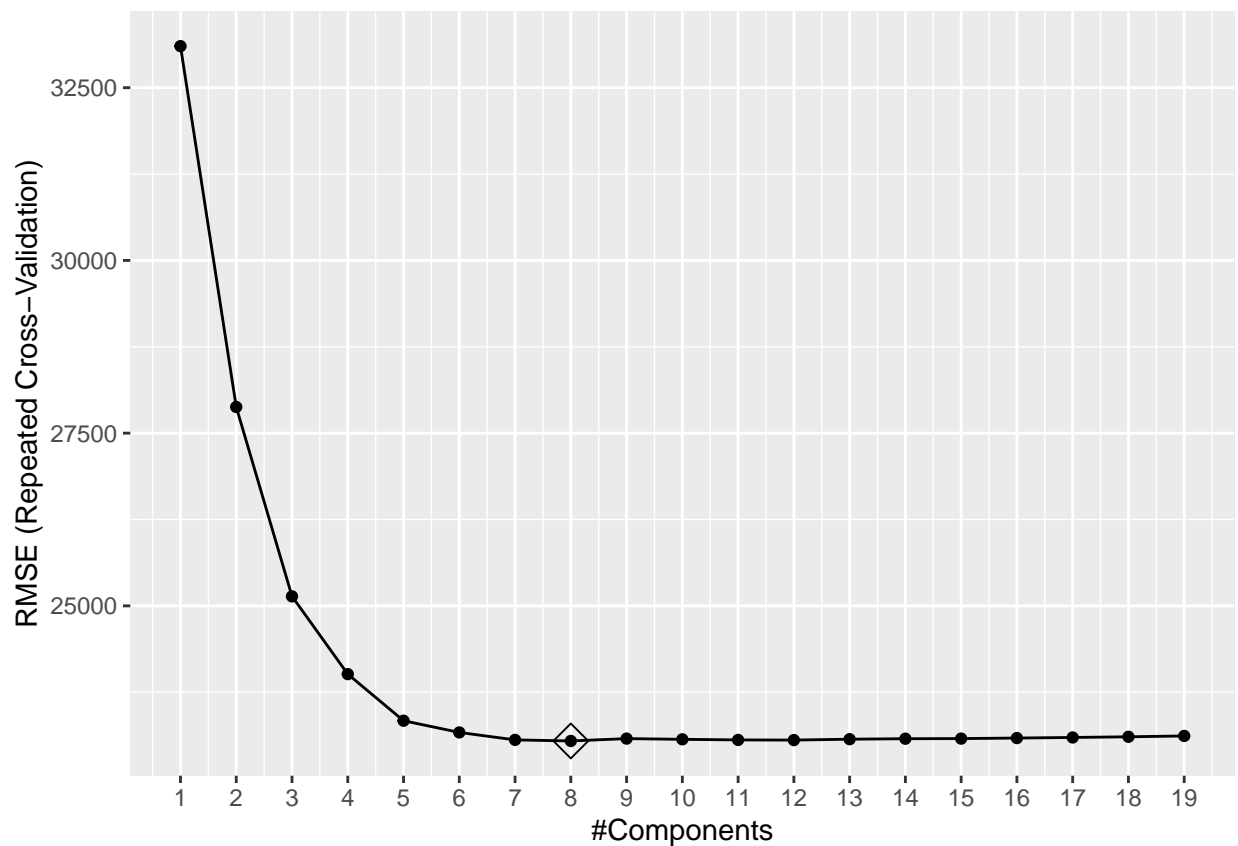
```
## [1] 440217938
```

8 components are included in this pls model which give the least cv rmsep. And the test error of this model is 4.4021794×10^8 .

PLS using caret

```
set.seed(1234)
pls.fit <- train(x, y,
  method = "pls",
  tuneGrid = data.frame(ncomp = 1:19),
  trControl = ctrl,
  preProcess = c("center", "scale"))

ggplot(pls.fit, highlight = TRUE) +
  scale_x_continuous(breaks = seq(0,20,1))#the number of components with the least RMSE is still 8
```



Model comparison

```
set.seed(1234)
resamp <- resamples(list(lm = fit_lm, lasso = lasso.fit, enet = fit_enet, pls = pls.fit))
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: lm, lasso, enet, pls
## Number of resamples: 50
##
## MAE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lm      14260.19 15738.77 16883.88 16773.69 17770.39 19630.86    0
## lasso   14196.65 15644.74 16779.82 16702.37 17715.29 19623.58    0
## enet    14104.43 15566.54 16726.15 16656.97 17695.97 19574.58    0
## pls     14110.75 15537.70 16743.88 16648.32 17586.87 19675.05    0
##
## RMSE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lm      20212.65 21478.04 22810.06 23132.49 24898.20 27059.42    0
## lasso   20184.81 21378.68 22792.62 23072.18 24511.54 27200.14    0
## enet    20107.10 21370.72 22829.14 23054.78 24487.70 27213.97    0
## pls     20135.01 21435.03 22804.56 23044.74 24593.52 27539.43    0
##
## Rsquared
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lm       0.8490164 0.8940725 0.9061873 0.9022587 0.9144933 0.9293063    0
## lasso    0.8488847 0.8946809 0.9066495 0.9026838 0.9146436 0.9287996    0
## enet     0.8500784 0.8946500 0.9066844 0.9028631 0.9151932 0.9284083    0
## pls      0.8491266 0.8937347 0.9053601 0.9029215 0.9159600 0.9292046    0

bwplot(resamp, metric = "RMSE")
```

